

# CSE 120

## PA4 Discussion

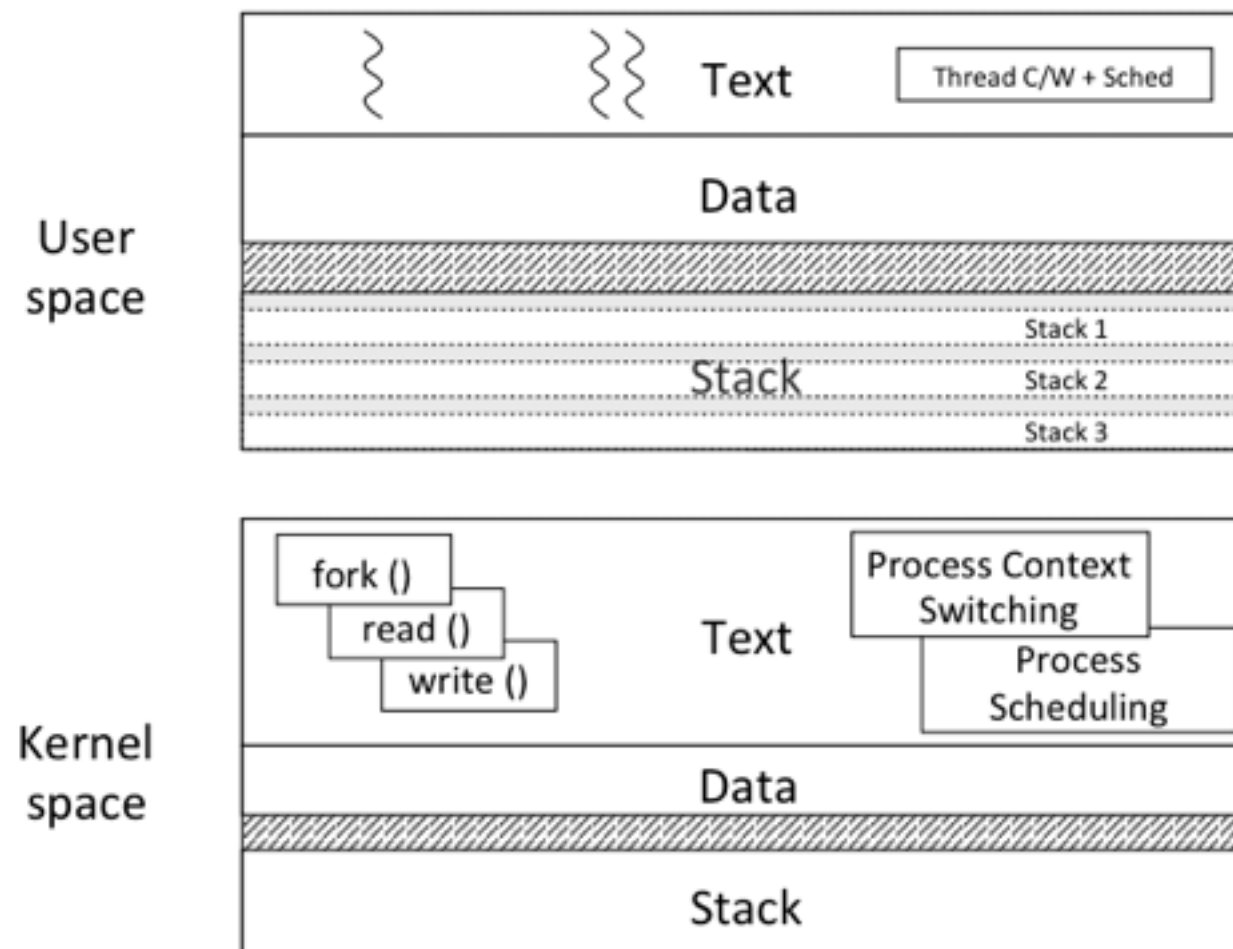
# PA4

- Implement a user-level thread package
- Entirely on user-level, no kernel modification
- You will be only turning in mythreads.c, which includes following functions:
  - MyInitThreads()
  - MySpawnThread(func, param)
  - MyYieldThread(t)
  - MyGetThread()
  - MySchedThread()
  - MyExitThread()

# Building Block

- `setjmp(env)` and `longjmp(env, t)` are the two utility function used to build thread packages
- `setjmp` and `longjmp` are very similar to `SaveContext` and `RestoreContext` in PA1.

# User-level Threads

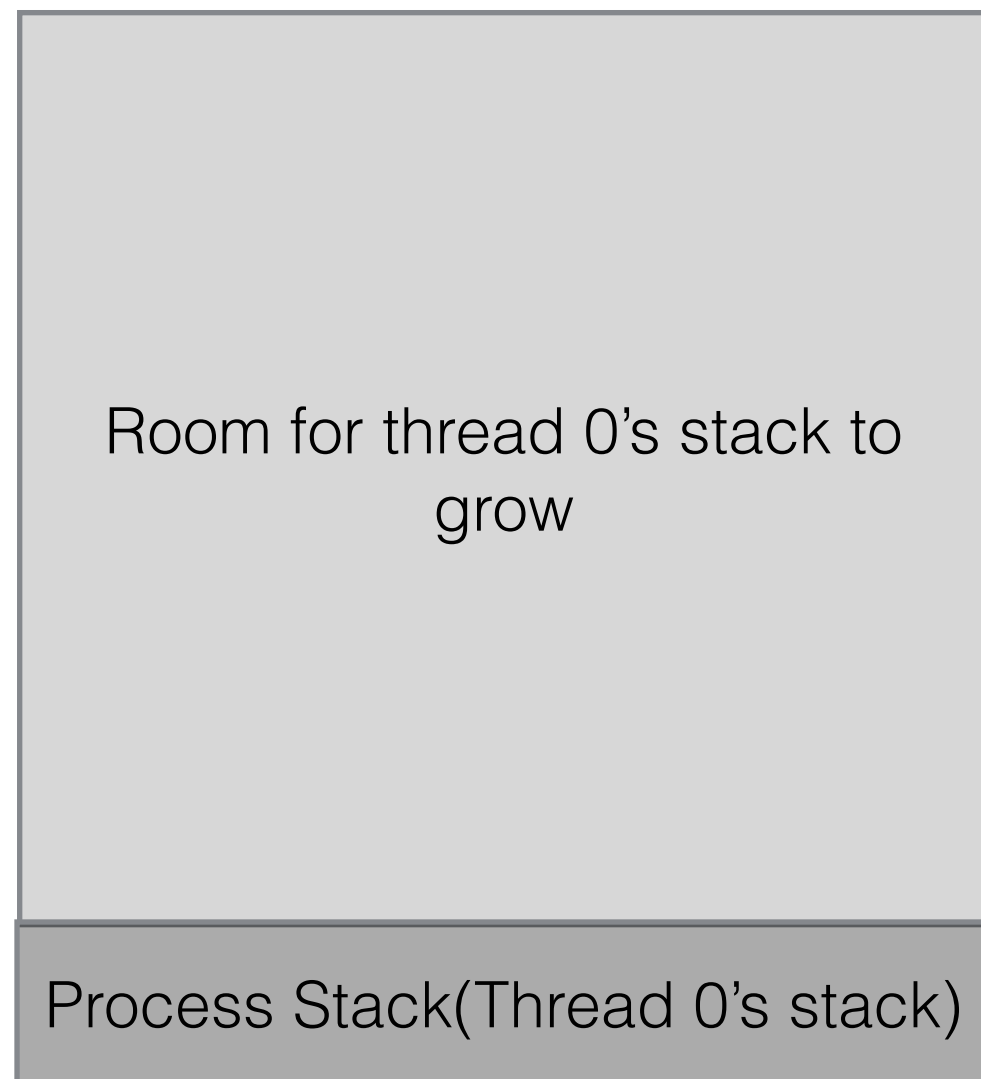


Stack is partitioned

# Stack Partition


- What is given

...



- What we desire

...



# InitThreads

- Partition the stack into MAXTHREADS parts
- Make use of dummy array to separate stacks of different threads
- Needs to generalize the code in `MySpawnThread(func, param)` to support partition into MAXTHREADS parts

# How to achieve partition

- Push the SP up by declaring a dummy array  
s[STACKSIZE]
- Save the execution environment (PC, SP, etc.)  
using setjmp(env)

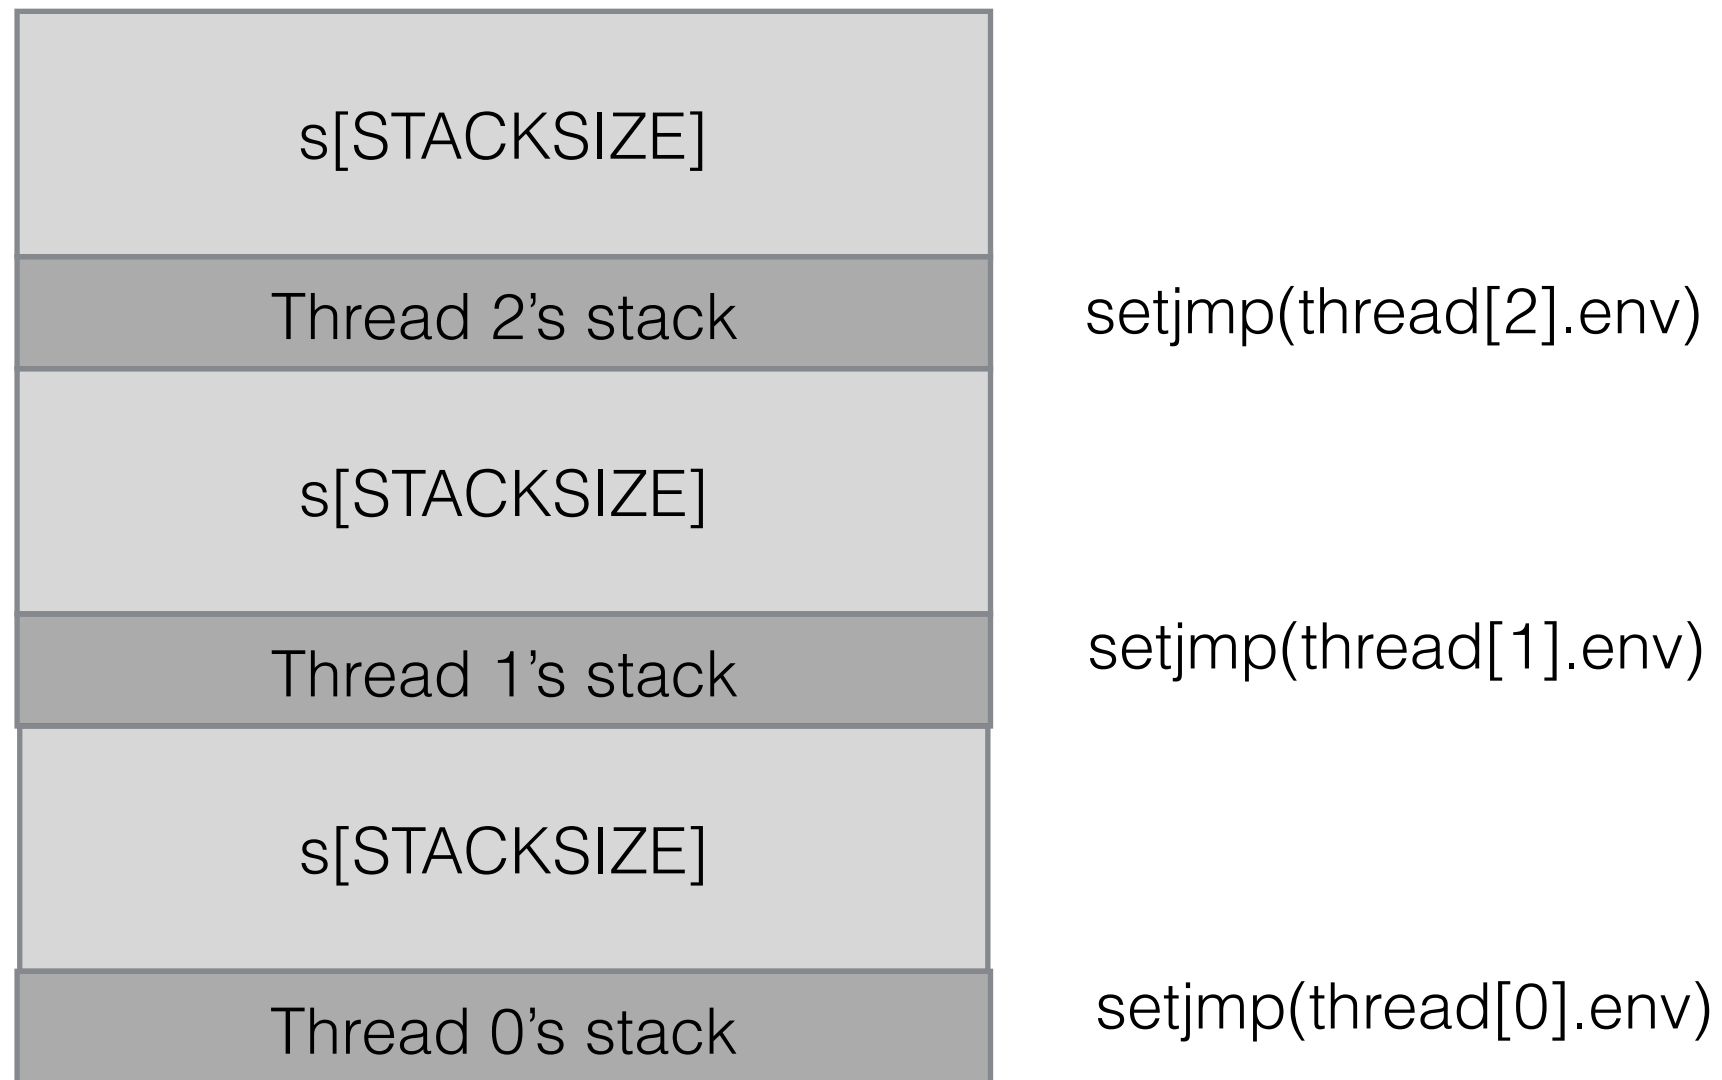
# Initial State





# After partition

...



# One Solution


```
for (int i=1; i<= MAXTHREADS; i++)  
{  
    char s[STACKSIZE]; //dummy array  
    Save the environment using setjmp(thread[i].env);  
}
```

Not correct. In each iteration of the for loop, `s[STACKSIZE-1]` points to the same location

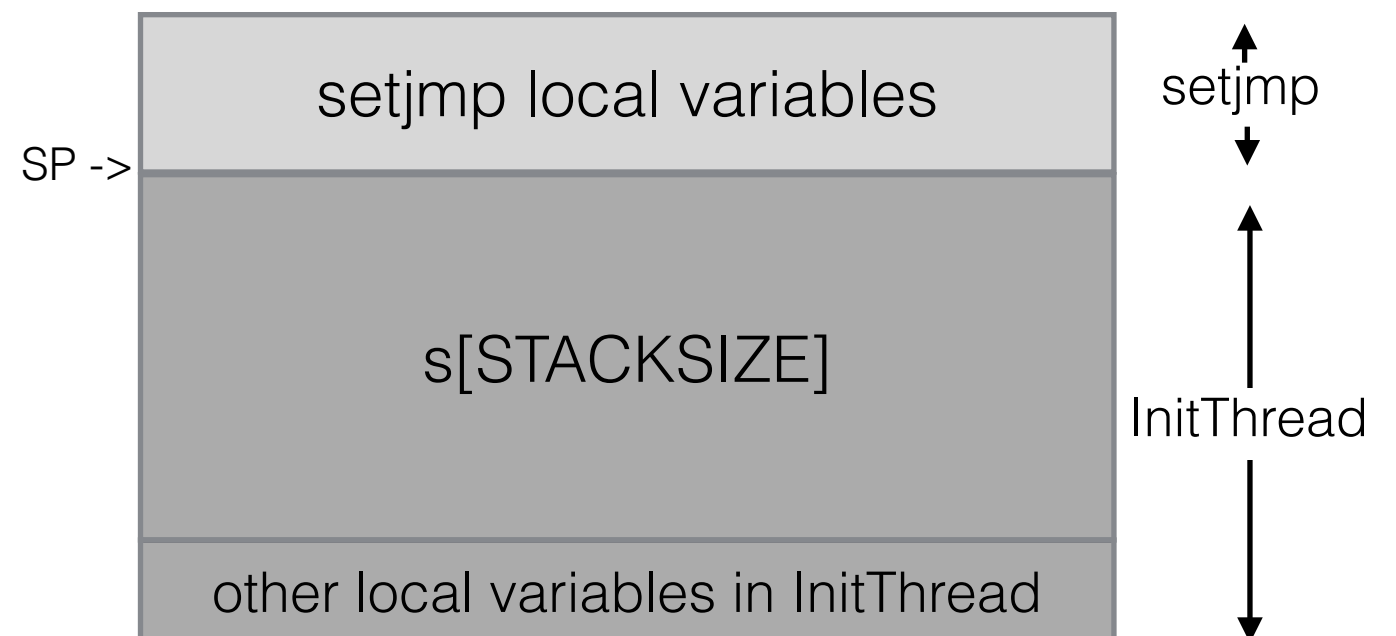
$i = 1$

- Memory snapshot when thread context is saved inside `setjmp(thread[i].env)`

Process perspective



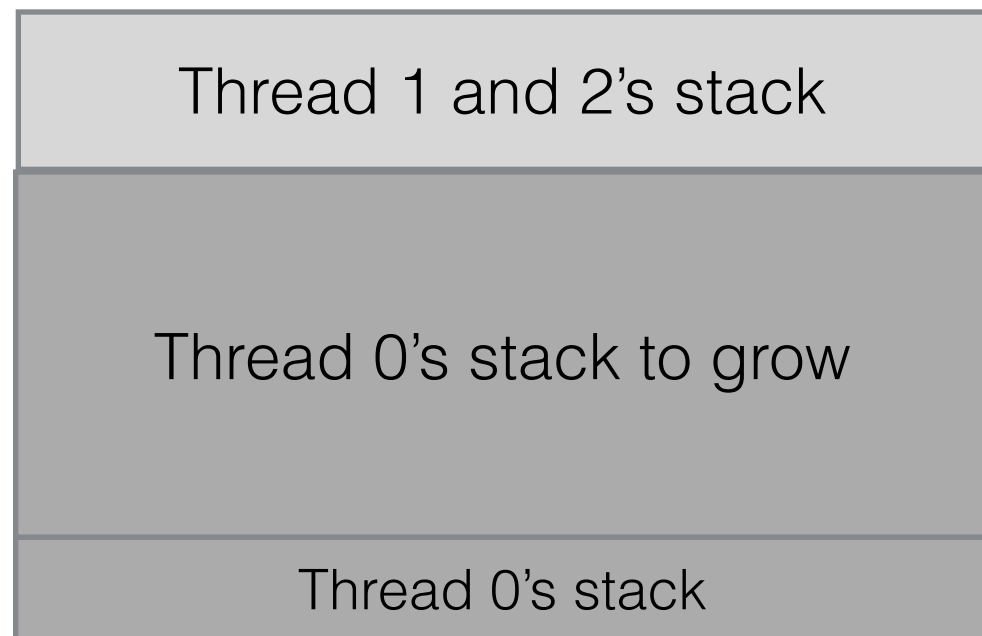
Kernel perspective



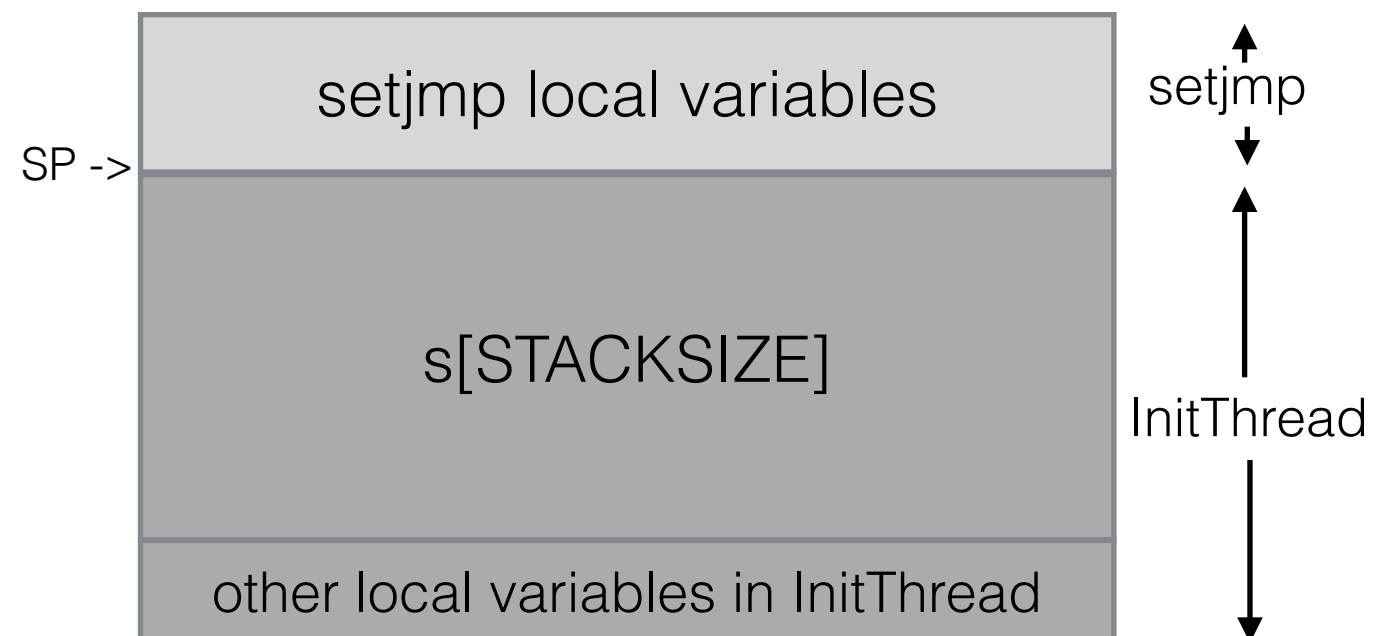
$i=2$

- Memory snapshot when thread context is saved inside `setjmp(thread[i].env)`

Process perspective



Kernel perspective



SP is not pushed up!

Thread 1 and 2's stack ended at the same memory location!

# How about this

```
for (int i=1; i<= MAXTHREADS; i++)
```

```
{
```

```
    char s[i*STACKSIZE]; //dummy array
```

```
    Save the environment using setjmp(thread[i].env);
```

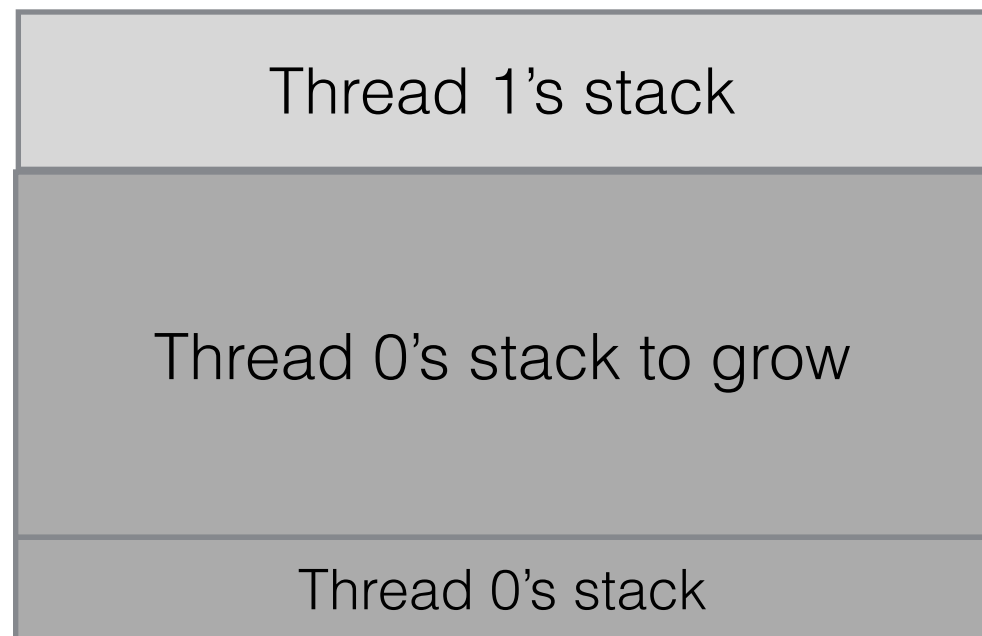
```
}
```

Now it works. In each iteration of the for loop, `s[i*STACKSIZE-1]` points to different location in different iterations

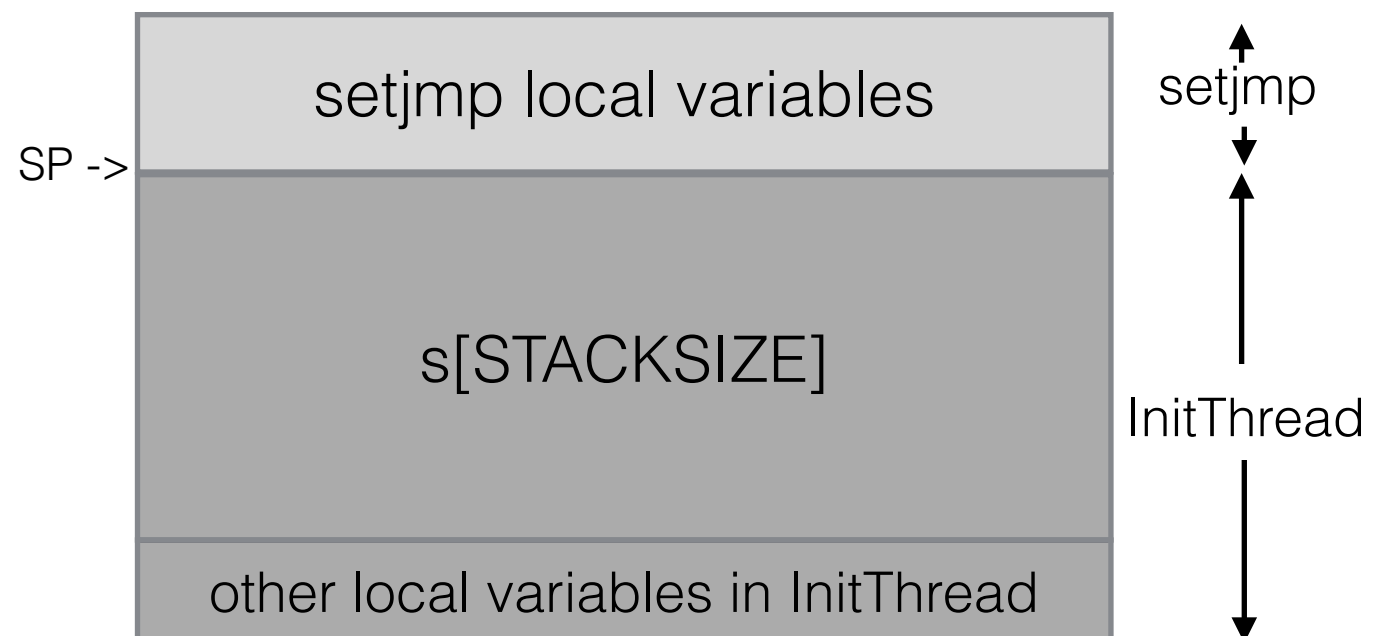
$i = 1$

- Memory snapshot when thread context is saved inside `setjmp(thread[i].env)`

Process perspective

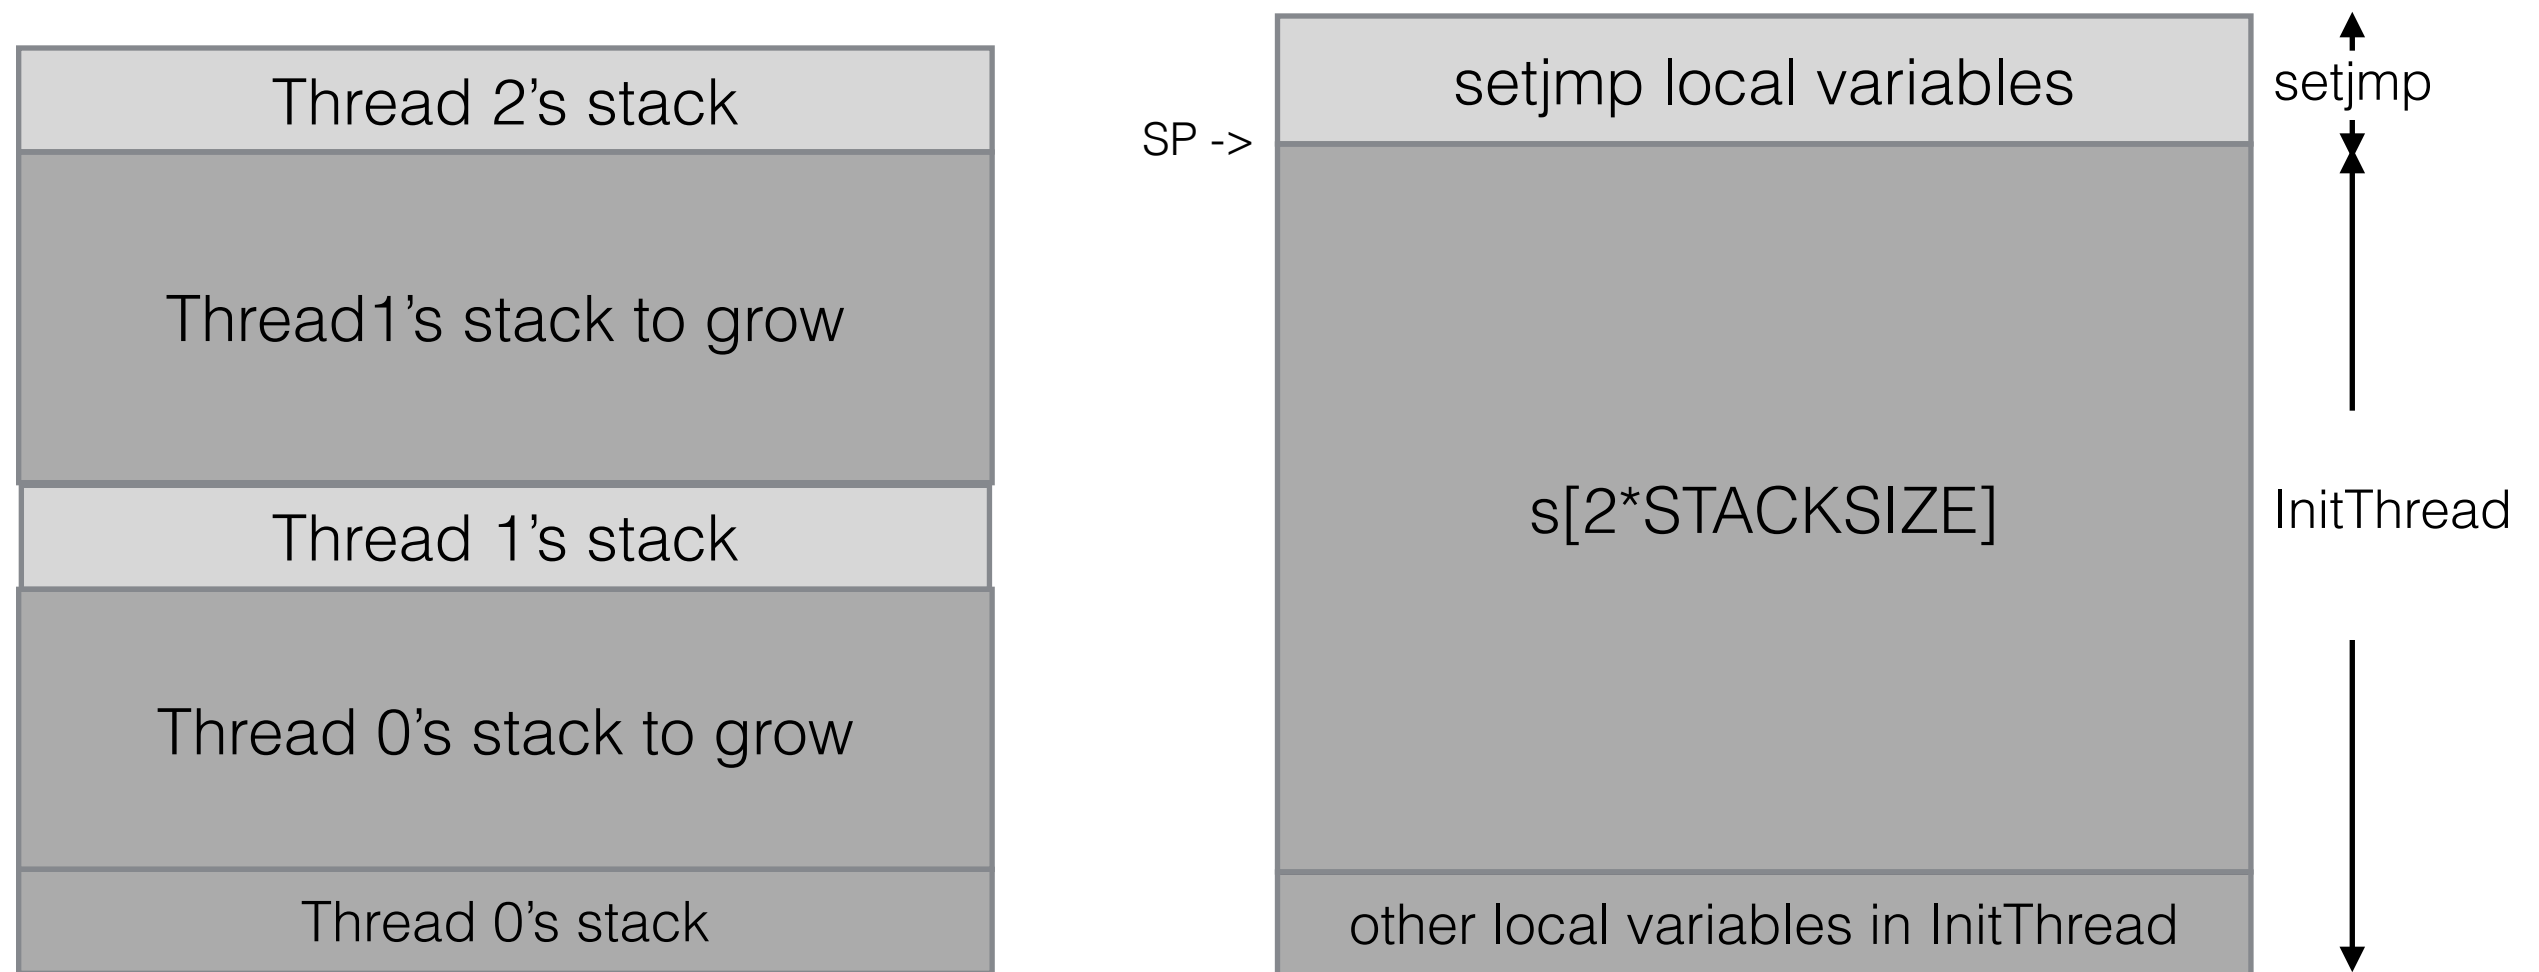


Kernel perspective



i=2

- Memory snapshot when thread context is saved inside `setjmp(thread[i].env)`



SP moves to a different location this time, so it would work!

# Alternative

- Using Recursion

```
stackPartition(int number_partitions) {
```

```
if (number_of_partitions <= 1) then return;
```

```
else {
```

```
char s[STACKSIZE];
```

```
Save the environment using setjmp(thread[i].env);
```

```
stackPartition(number_partitions-1);
```


```
}
```

```
}
```



# Example

- Partition the stack into 5 parts, initially calls `stackPartition(5)`.




# Example

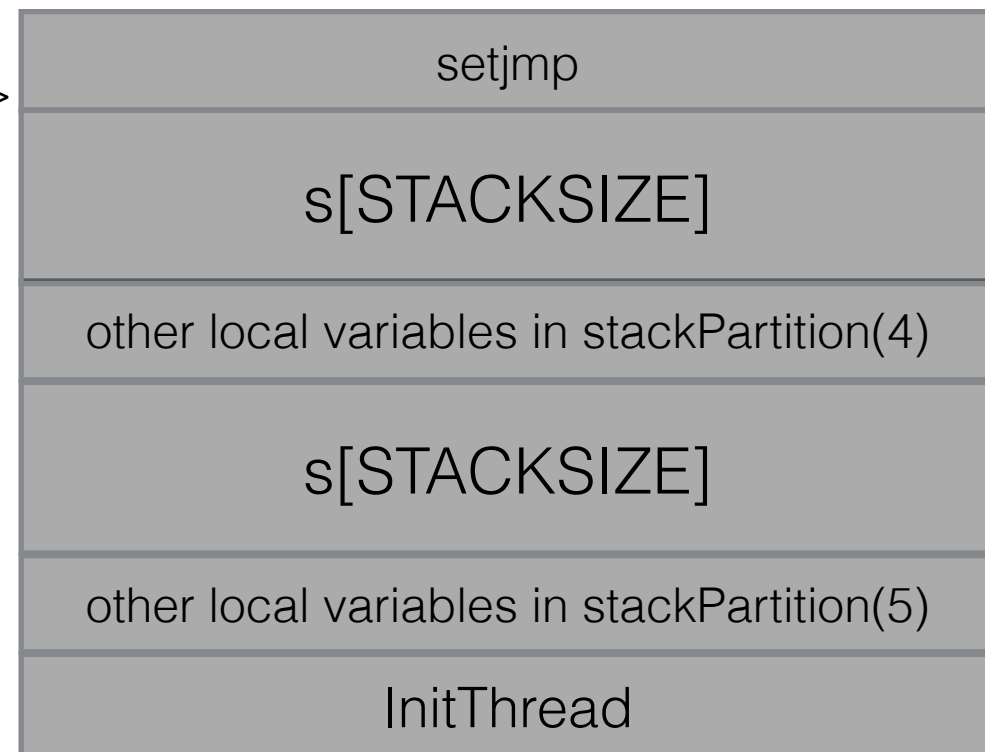
- Inside `stackPartition(5)`, we call `stackPartition(4)`

...

...



SP ->



# MySpawnThread

- `MySpawnThread(func, param)` spawns a new thread which executes `func(param)`
- Returns thread-id of the thread spawned.
- Incremental assignment of thread id.

# Thread Id Example

Thread Id



0	T0
1	T1
2	T2
3	T3
4	T4
5	T5
6	T6
7	
8	
9	

Thread 0 ~ 6  
spawned

# Thread Id Example

0	T0
1	T1
2	
3	T3
4	T4
5	
6	T6
7	
8	
9	

Thread 2, 5  
exits

# Thread Id Example

0	T0
1	T1
2	
3	T3
4	T4
5	
6	T6
7	T7
8	
9	

Thread 7 spawned  
Where should be put thread 7?

# Thread Id Example

0	T0
1	T1
2	
3	T3
4	T4
5	
6	T6
7	T7
8	T8
9	T9

Thread 8, 9 spawned

# Thread Id Example

0	T0
1	T1
2	T10
3	T3
4	T4
5	
6	T6
7	T7
8	T8
9	T9

Thread 10 spawned



# Thread Id Example

0	
1	T1
2	T10
3	
4	T4
5	
6	T6
7	T7
8	T8
9	T9

Thread 0, 3 exits

# Thread Id Example

Thread 11 spawned  
Where should we put thread 11?

0	
1	T1
2	T10
3	T11
4	T4
5	
6	T6
7	T7
8	T8
9	T9

← Last created

# Thread Id Example

0	T13
1	T1
2	T10
3	T11
4	T4
5	T12
6	T6
7	T7
8	T8
9	T9

Thread 12, 13 spawned

# MySpawnThread

- Keep track of the last assigned thread id
- When need to spawn a thread, find the next available id after the last assigned thread id.
- Also need to store the function pointer `void (*f)()`, `int p`, which the spawned thread would need in the future
- Needs to add some variables to the thread table

# MyYieldThread

- MyYieldThread(t) causing the currently running thread to yield to thread t
- Returns the id of the thread which called MyYieldThread
- Very similar to MyContextSwitch in PA1

# Example

Thread 1 starts first

Thread 1:

```
do some work  
x = MyYieldThread(2)  
// x = 2  
continue working
```

Thread 2:

```
do some work  
y = MyYieldThread(1)  
// y = 1  
continue working
```

# MyYieldThread

- Use `setjmp` to save the context of current thread
- Use `longjmp` to give up CPU to the specified thread.
- We can set the second parameter of `longjmp(env, t)` to be the currently running thread id. So after `longjmp` returns, the thread that gets control of CPU will know what the previous running thread is (the id to return in `MyYieldThread`).

# Edge Cases

- Yielding to self
- Yielding to an invalid thread id
  - MyYieldThread(-3)
  - MyYieldThread(MAXTHREAD + 10)



# MySchedThread

- Pick one thread to run according to the scheduling policy
- In PA4, we are implementing a FIFO policy
- Once you determined which thread to run next, you can just make use of MyYieldThread

# FIFO queue

- First in first out, last in last out
- Things become tricky when MyYieldThread comes into play and change the location of interior elements
- MySpawnThread would add new elements into the queue

# Scheduling Example

Running Thread: 1

Queue: [2,0,3,4]  
          ↑  
          head

Thread 1 calls MySchedThread(), which thread should run next?

Where should we put thread 1 in the queue?

# Scheduling Example

Running Thread: 2

Queue: [0,3,4,1]

↑  
head

# Scheduling Example

Running Thread: 2

Queue: [0,3,4,1]  
          ↑  
          head

Thread 2 calls MySpawnThread(), creates thread 5, where should we put thread 5?

# Scheduling Example

Running Thread: 2

Queue: [0,3,4,1]  
          ↑  
          head

Thread 2 calls MySpawnThread(), creates thread 5, where should we put thread 5?

# Scheduling Example

Running Thread: 2

Queue: [0,3,4,1,5]

↑  
head

# Scheduling Example

Running Thread: 2

Queue: [0,3,4,1,5]  
          ↑  
          head

Thread 2 calls MyYieldThread(4), what should the queue look like?



# Scheduling Example

Running Thread: 4

Queue: [0,3,1,5,2]

↑  
head

# MyExitThread

- The function to call when thread is done with its job
- Always call MyExitThread() after the function execution line in mythreads.c
- It releases the slot in thread table so threads spawned later on can make use of that.
- Needs to reset the **env** to initial state (you may need to add some variables inside thread table).

# PA4 Checklist

- Pre-partition stack space in MyInitThread
- MySpawnThread chooses the right thread Id
- MySchedThread picks the right thread to run
- MyYieldThread should switch to specified thread and return the id of the thread that gave up CPU.
- MyExitThread reset env so thread slot can be reused.